

Intro to Social Science Data Analysis

Lecture 6: Data Visualisation in R

Christopher Gandrud

October 7, 2012

- 1 Assignment 2
- 2 Recap
- 3 Principles of Graphics Excellence
- 4 Base R Graphics
- 5 ggplot2
- 6 googleVis

Assignment 2

Due: Friday 19 October

Describe at least **3** variables in a data set.

You need to select a **range of descriptive statistical tools**. The tools should include both **numerical descriptive statistics** and **graphics**.

These tools should describe the variables':

- ▶ central tendency,
- ▶ variation,
- ▶ their relationships with the other variables.

The descriptions need to be discussed **in paragraph form**.

The description must be **reproducible**. So you should email me the link to a Dropbox folder with:

- ▶ the .csv data set,
- ▶ the .Rmd R markdown file,
- ▶ the final .html file.

Quick Quiz (1)

When you describe data, what **two** things do you always need to discuss?

Why do you need to describe both things?

Give examples for data at different measurement levels.

Quick Quiz (2)

What is the difference between the **population** mean and the **sampling** mean?

Why would you log transform a variable?

Today

Last week: we largely learned how to describe our data *numerically*.

Today: we will learn how to present our data with *graphics*.

We will learn both how to create graphics in R, but also the principles of effective statistical graphics.

Many of the things we learn today will also apply to inferential statistics.

The first part of this lecture is based on Tufte (2001)

Many of the examples are from the Junk Charts Blog (<http://junkcharts.typepad.com/>).

We will also use the World Bank data we downloaded last class.

R Source Code at: <http://bit.ly/0TWEGS>

Why Graphics?


Why use graphics? Why not just describe all of our data in tables?

```
# Create data frame with GDP/Capita & Infant Mort.  
DataDump <- InfantNoMiss[,  
  c("GDPperCapita", "InfantMortality")]
```

```
# Show data
```

```
DataDump
```

##	GDPperCapita	InfantMortality
## 7	38959.8	6.3
## 8	425.1	76.2
## 9	13829.8	7.2
## 10	3795.7	14.1
## 11	2803.3	17.2
## 12	4068.5	100.1
## 13	7665.1	13.4
## 15	45638.1	3.6
## 16	42101.4	4.3
## 18	4950.3	41.9
## 19	4534.1	7.1
## 20	13181.3	16.9



Goal of Statistical Graphics:

The efficient communication of complex
quantitative ideas.

Tufte's Principles for Excellent Statistical Graphics (2001, 13) (a selection):

- ▶ show the data
- ▶ encourage the eye to compare differences in the data
- ▶ serve a clear purpose
- ▶ avoid distorting the data
- ▶ be closely integrated with the text

Tufte's Principles for Excellent Statistical Graphics (2001, 13) (a selection):

- ▶ show the data
- ▶ encourage the eye to compare differences in the data
- ▶ serve a clear purpose
- ▶ avoid distorting the data
- ▶ be closely integrated with the text

Tufte's Principles for Excellent Statistical Graphics (2001, 13) (a selection):

- ▶ show the data
- ▶ encourage the eye to compare differences in the data
- ▶ serve a clear purpose
- ▶ avoid distorting the data
- ▶ be closely integrated with the text

Tufte's Principles for Excellent Statistical Graphics (2001, 13) (a selection):

- ▶ show the data
- ▶ encourage the eye to compare differences in the data
- ▶ serve a clear purpose
- ▶ avoid distorting the data
- ▶ be closely integrated with the text

Tufte's Principles for Excellent Statistical Graphics (2001, 13) (a selection):

- ▶ show the data
- ▶ encourage the eye to compare differences in the data
- ▶ serve a clear purpose
- ▶ avoid distorting the data
- ▶ be closely integrated with the text

Show the Data

Encourage the eye to compare differences.

Serve a purpose.

Show the data

Show the data, not other things like silly graphics, or unnecessary words.

Have high **data ink** ratio.


$$\text{Data Ink Ratio} = \frac{\text{data} - \text{ink}}{\text{total ink}} \quad (1)$$

Example


Encourage the Eye to Compare Differences

How did the budgets change?


(Orange is 2013, Blue is 2012)



A Little Better



Percentage Change in Departmental Spending 2012 to 2013



Avoid distorting the data.

Special case: Circles.

Avoid Circles! (1)

In general: Avoid using the *size* of a circle to mean something!

This includes avoiding:

- ▶ Bubble charts
- ▶ Pie charts

Avoid Circles! (1)

In general: Avoid using the *size* of a circle to mean something!

This includes avoiding:

- ▶ Bubble charts
- ▶ Pie charts

Avoid Circles! (2)

Why?

Circles can distort data.

- ▶ It is very difficult to compare their size
- ▶ The Ebbinghause Illusion!

Avoid Circles! (2)

Why?

Circles can distort data.

- ▶ It is very difficult to compare their size
- ▶ The Ebbinghause Illusion!

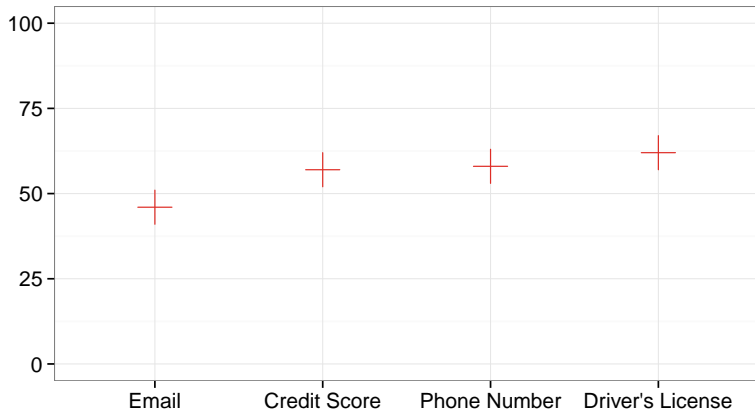
Order the 4 circles from largest to smallest.

The circles are on a scale of 0-100, so how much bigger are each of the circles relative to each other?



Order the 4 bars from largest to smallest.

How much bigger are each of the circles relative to each other?




The circles basically tell you nothing that a simple table could do better.



Ebbinghaus Illusion

Which circle is bigger?



Colours and Data Distortions

Special Case 2.

Colour & Context Illusions



Only give graphical features (e.g. bars in a bar chart) different colours if it **means something** in the data.

See: Stephen Few

http://www.perceptualedge.com/articles/visual_business_intelligence/rules_for_using_color.pdf

Colours should be used to:

- ▶ highlight particular data,
- ▶ group items,
- ▶ encode quantitative values.


Colours should be used to:

- ▶ highlight particular data,
- ▶ group items,
- ▶ encode quantitative values.

Colours should be used to:

- ▶ highlight particular data,
- ▶ group items,
- ▶ encode quantitative values.

Bad



Good



Colour blindness.

Be colour blind friendly.

Colour blindness:

People who are colour blind can have difficulty distinguishing between red-green and blue-yellow.

About 5-8% of men are colour blind.

We need to choose colour schemes for our graphics that are **colour blind friendly**.

For more information see <http://www.usability.gov/articles/newsletter/pubs/022010new.html>.

Remember:

Graphics are only as good as what you put in them.

Silly data and statistics will always create silly graphs.

Base R Graphics

Last week we saw that R has some basic graphics functions like:

- ▶ `plot`


- ▶ `histogram`

Last week we saw that R has some basic graphics functions like:


- ▶ `plot`
- ▶ `histogram`

A Basic Scatter Plot

```
# Create a basic scatter plot  
with(MortalityGDP,  
      plot(x = InfantMortality, y = GDPperCapita))
```



```
# Basic scatter: axis labels & rescale GDP/Capita
with(MortalityGDP,
  plot(x = InfantMortality,
        y = (GDPperCapita/1000),
        xlab = "Infant Mortality per 1000 Births",
        ylab = "GDP/Capita (1,000s of US$)",
        main =
          "Scatterplot of\n GDP/Capita & Infant Mort."))
```



Graphics with ggplot2

Why ggplot2?

The ggplot2 package

The ggplot2 package allows us to do much more than base R graphics.

What is ggplot2?

“gg” means “Grammar of Graphics”

(2 just means it is the second one)


Each plot is made of **layers**. Layers include the coordinate system (x-y), points, labels, etc.

Each layer has **aesthetics** (aes) including the x & y, size, shape, and colour.

The main layer types are called **geometrics** (geom). These include lines, points, density plots, bars, and text.

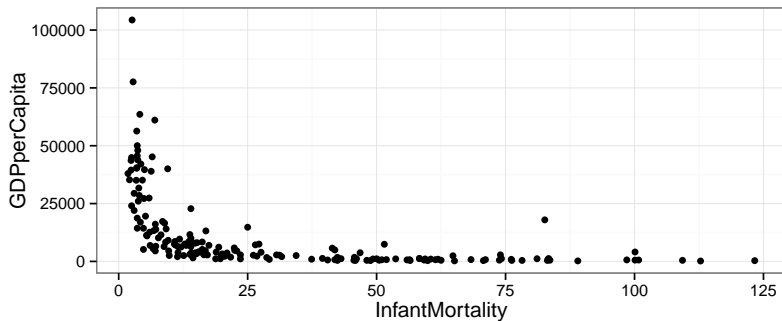
Simple Example

```
# Scatterplot  
ggplot(data = MortalityGDP, aes(x = InfantMortality,  
                                y = GDPperCapita)) +  
  geom_point()
```



Simple Example with Blank Theme

```
# Scatterplot
ggplot(data = MortalityGDP, aes(x = InfantMortality,
                                y = GDPperCapita)) +
  geom_point() +
  theme_bw(base_size = 13)
```

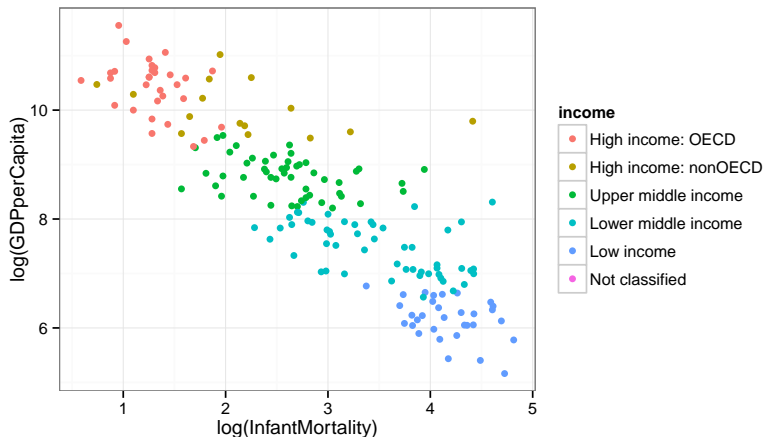


Colours

There are a number of ways to specify colours in `ggplot2`.

The simplest way is to let `ggplot` choose the colours for you.

```
# Create scatter plot divided by region
ggplot(data = InfantNoMiss, aes(log(InfantMortality),
                                log(GDPperCapita))) +
  geom_point(aes(colour = income)) +
  theme_bw()
```



Colors! (2)

There are many ways to pick specific colors.

In this class we will mainly use **hexadecimal** colours.

This is probably the most commonly used system for choosing colours on the web.

Every colour is given six digits.

A good website for getting hexadecimal colour schemes is:
<http://colorbrewer2.org/>.


```
# Create colour vector
```

```
Colours <- c("#1B9E77", "#D95F02", "#7570B3",  
             "#E7298A", "#66A61E", "#E6AB02")
```


```
# Create graph
```

```
ColourScatter <- ggplot(data = InfantNoMiss,  
                        aes(log(InfantMortality),  
                           log(GDPperCapita))) +  
  geom_point(aes(colour = income)) +  
  scale_color_manual(values = Colours) +  
  xlab("\nLog Infant Mortality") +  
  ylab("Log GDP/Capita\n") +  
  ggtitle("Log Transformed Data\n") +  
  theme_bw()
```

```
# Show scatter plot
```

```
ColourScatter
```

Log Transformed Data



```
# Create a violin Plot
Violin <- ggplot(InfantNoMiss, aes(
  factor(DumMort),
  log(GDPperCapita))) +
  geom_violin(fill = "#E7298A",
    colour = "#E7298A",
    alpha = I(0.5)) +
  geom_jitter(color = "#7570B3") +
  xlab("\n Infant Mortality") +
  ylab("Log GDP.Capital\n") +
  theme_bw(base_size = 16)
```

```
# Create a violin Plot
```

```
Violin
```




More Information:

`http://docs.ggplot2.org/current/index.html`

Maps with googleVis

Maps with ggplot2

Figure: RStudio Downloads (1 Month in 2012)



Source: <http://blog.rstudio.org/2012/10/01/where-in-the-world/#comments>

googleVis

We can create HTML based graphics with GoogleVis.

HTML graphics can be **interactive**.

This includes **motion charts & maps**.

Map Example

More Information:

Class Example (all files): <http://bit.ly/WdeMAw>

googleVis Info:

<http://code.google.com/p/google-motion-charts-with-r/>

Many people use R to create professional graphics.

For example: see the New York Times' graphics blog: <http://chartsnthings.tumblr.com/>

They often use R in combination with Adobe Illustrator.

See Nathan Yau's Book *Visualize This: The Flowing Data Guide to Design, Visualization, and Statistics* (<http://book.flowingdata.com/>).

References I

Tufte, Edward R. 2001. The Visual Display of Quantitative Information. Cheshire, Connecticut: Graphics Press.